

<Arnav Mehta>

<arnavm7>

<AL>

**Name:**

**NetID:**

**Section:**

## ECE 408/CS483 Milestone 3 Report

- List Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images from your basic forward convolution kernel in milestone 2. This will act as your baseline this milestone.

Batch Size
Op Time 1
Op Time 2
Total Execution Time
Accuracy
100
<0.17832 ms>
<0.636396 ms>
<1.293 s>
<0.86>
1000
<1.60053 ms>
<6.11565 ms>
<10.143 s>
<0.886>
10000
<15.8478 ms>
<61.2404 ms>
<1m39.470 s>
<0.8714>

- **Optimization 1: <Using Streams to Overlap Computation with Data Transfer>**

- Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (**2 points**)
- ☐ Shared memory matrix multiplication and input matrix unrolling (**3 points**)
- ☐ Kernel fusion for unrolling and matrix-multiplication (**2 points**)
- ☐ Weight matrix in constant memory (**1 point**)
- ☐ Tuning with restrict and loop unrolling (**3 points**)

- ☐ Sweeping various parameters to find best values (**1 point**)
- ☐ Multiple kernel implementations for different layer sizes (**1 point**)
- ☐ Input channel reduction: tree (**3 point**)
- ☐ Input channel reduction: atomics (**2 point**)
- ☐ Fixed point (FP16) arithmetic. (**4 points**)
- ☐ Using Streams to overlap computation with data transfer (**4 points**)

I CHOSE TO IMPLEMENT THE STREAMS

#### OPTIMIZATION

- ☐ An advanced matrix multiplication algorithm (**5 points**)
- ☐ Using Tensor Cores to speed up matrix multiplication (**5 points**)
- ☐ Overlap-Add method for FFT-based convolution (**8 points**)
- ☐ Other optimizations: please explain

<answer here>

- How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

<This optimization works by overlapping tasks to reduce the bottleneck of data transfer. Unlike the traditional flow of first sending data to the GPU, having the kernel perform computations, and then sending it back, streams allow for these actions to happen concurrently. A stream is a queue of actions that keeps track of dependencies.

As a result of streams input data can be split up, and some can be sent to the GPU while the GPU's computation on another piece of data gets sent back to the host. In this optimization I had 10 different streams, with each stream processing a batch size of 10 images. In a loop structure I first had the streams send data to the GPU, which are independent activities. These go to the data transfer engine/driver. Then, the streams all launched the kernel via the kernel engine. Then, at the end of an iteration, the streams got the data back from the GPU.

There is an issue where subsequent transfers from host to device are hindered by the queue and the dependency of transfers from the device to the host on the kernel execution. I tried to increase the number of streams so that this would not be a problem at the beginning, which seemed like it could reduce overhead for small batch sizes. I do not think this optimization would synergies particularly well with some other operations, like constant memory for example, due to the large number of kernel calls and thus loading from global memory to constant memory.>

- List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size
Op Time 1
Op Time 2
Total Execution Time
Accuracy
100
<0.001097 ms>
<0.001024 ms>
<0m1.181s>
<0.86>
1000
<0.001076 ms>
<0.001056 ms>
<0m9.747 s>
<0.886>
10000
<0.001166 ms>
<0.001587 ms>
<1m37.128 s>
<0.8714>

- Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of  
 <This implementation was partially successful in improving performance. Total execution times were less than the baseline (for example in 1000 batch case the streams had a total time of 9.747 seconds which was less than 10.143 seconds in the baseline case. The op times here do not accurately measure the speed because memcpy is happening so frequently. The profiling is as follows:

Generating CUDA API Statistics...  
 CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum
Name					
-----	-----	-----	-----	-----	-----
85.3	1254657823	4000	313664.5	29993	1192353
cudaMemcpyAsync					
12.6	185983604	8	23247950.5	73613	182909037
cudaMalloc					
1.6	23948217	2004	11950.2	3123	12175404
cudaLaunchKernel					
0.2	3209323	4	802330.7	10016	3160489
cudaMemcpy					
0.2	2891419	8	361427.4	68036	986158
cudaFree					

0.0	292814	4	73203.5	21035	123284
cudaHostRegister					
0.0	145866	20	7293.3	1394	86917
cudaStreamCreate					
0.0	121925	20	6096.2	1831	38486
cudaStreamDestroy					
0.0	39828	8	4978.5	1948	8350
cudaDeviceSynchronize					

Generating CUDA Kernel Statistics...

Generating CUDA Memory Operation Statistics...

CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum
Name					
100.0	134937686	2000	67468.8	22912	124767
conv_forward_kernel					
0.0	2816	2	1408.0	1376	1440
do_not_remove_this_kernel					
0.0	2720	2	1360.0	1312	1408
prefn_marker_kernel					

CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum
Name					
73.2	136441620	2000	68220.8	57056	91104 [CUDA memcpy DtoH]
26.8	49922461	2004	24911.4	1536	27967 [CUDA memcpy HtoD]

CUDA Memory Operation Statistics (KiB)

Total	Operations	Average	Minimum	Maximum
Name				

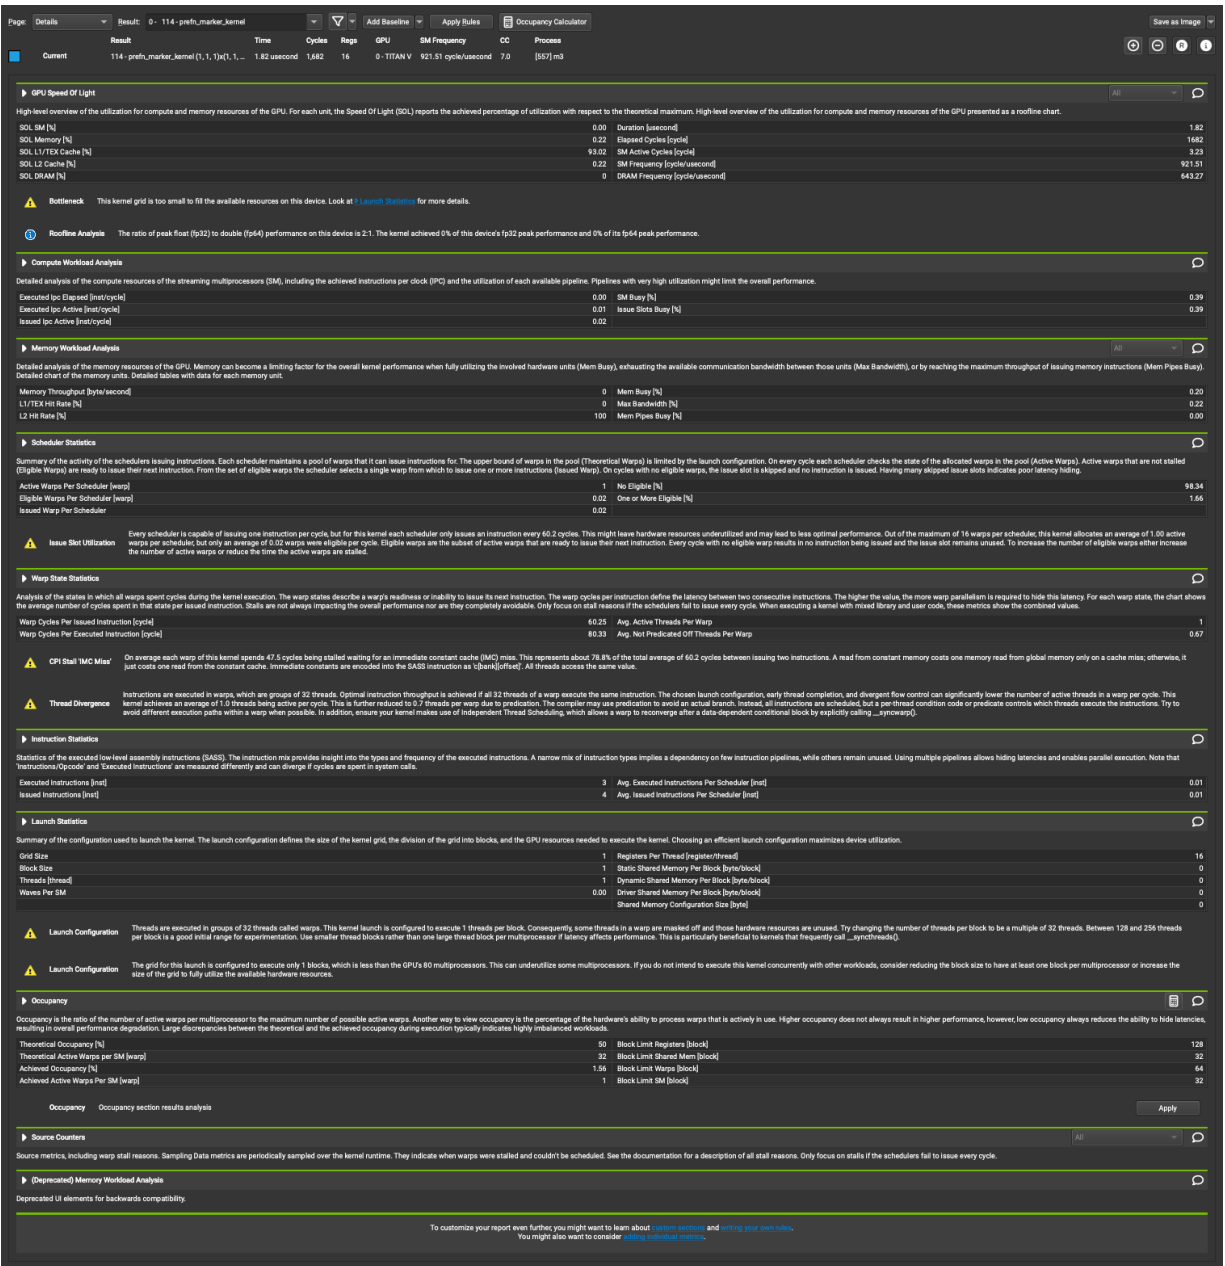
1722500.0	2000	861.0	722.500	1000.0 [CUDA memcpy DtoH]
538919.0	2004	268.9	0.004	288.0 [CUDA memcpy HtoD]

Generating Operating System Runtime API Statistics...  
Operating System Runtime API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum
Name					
33.3	96197588197	976	98563102.7	23513	100413091
sem_timedwait					
33.3	96176948240	975	98643023.8	35416	100553552
poll					
22.2	64147271031	2	32073635515.5	23453954467	40693316564
pthread_cond_wait					
11.1	32008316231	64	500129941.1	500049524	500182970
pthread_cond_timedwait					
0.0	92081956	923	99763.8	1036	18818041
ioctl					
0.0	22738425	9072	2506.4	1270	37548
read					
0.0	17088224	26	657239.4	1532	17015224
fopen					
0.0	3636000	98	37102.0	1178	1591611
mmap					
0.0	1137756	101	11264.9	4138	26753
open64					
0.0	466864	1	466864.0	466864	466864
pthread_mutex_lock					
0.0	308047	5	61609.4	40480	76885
pthread_create					
0.0	134849	3	44949.7	41962	48999
fgets					
0.0	84047	19	4423.5	1374	14813
munmap					
0.0	83913	3	27971.0	3896	55330
fopen64					
0.0	64916	15	4327.7	2182	7925
write					
0.0	56669	8	7083.6	2311	9958
fflush					

0.0	30387	5	6077.4	2845	7864
open					
0.0	26328	10	2632.8	1014	10360
fclose					
0.0	14297	2	7148.5	4631	9666
pthread_cond_signal					
0.0	13538	2	6769.0	5924	7614
socket					
0.0	12296	8	1537.0	1216	1865
fcntl					
0.0	7903	1	7903.0	7903	7903
pipe2					
0.0	7613	1	7613.0	7613	7613
connect					
0.0	6270	3	2090.0	1057	3556
fwrite					
0.0	2224	1	2224.0	2224	2224
bind					

>



What references did you use when implementing this technique?  
<I used course lecture 22 and the CUDA API programming guide for syntax for the streams.>

Please Paste your kernel code for this optimization. Your code

should include the non-trivial code that you have changed for this optimization. For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling

```
<cudaStream_t cudaArr[10];
for(int i = 0; i < 10; i++){
    cudaStreamCreate(&cudaArr[i]);
}

int inter_batch = Channel * Height * Width;
int inter_out = Map_out * (Height - K + 1) * (Width
- K + 1);
float * temp = (float*) *device_output_ptr;
float * temp2 = (float*) host_output;

    cudaMemcpy(*device_mask_ptr, host_mask,
sizeof(float) * Channel * Map_out * K * K,
cudaMemcpyHostToDevice);

    cudaHostRegister(&temp2, Batch * inter_out,
cudaHostRegisterDefault);
    cudaHostRegister(&host_input, Batch * inter_batch,
cudaHostRegisterDefault);
    for(int i = 0; i < Batch; i+= seg_size * 10){

        for(int s = 0; s < 10; s++){
            cudaMemcpyAsync(*device_input_ptr + (i + s *
seg_size) * inter_batch, host_input + (i + s *
seg_size) * inter_batch, seg_size * sizeof(float) *
inter_batch, cudaMemcpyHostToDevice, cudaArr[s]);
        }

        for(int s = 0; s < 10; s++){
            conv_forward_kernel<<<dimGrid, dimBlock, 0,
cudaArr[s]>>>(*device_output_ptr + (i + s * seg_size) *
inter_out,*device_input_ptr + (i + s * seg_size) *
inter_batch, *device_mask_ptr, seg_size, Map_out,
Channel, Height, Width, K);
        }
    }
```



```

    for(int s = 0; s < 10; s++){
        cudaMemcpyAsync(temp2 + (i + s * seg_size) *
inter_out, temp + (i + s * seg_size) * inter_out,
seg_size * inter_out * sizeof(float),
cudaMemcpyDeviceToHost, cudaArr[s]);
    }

}

```

>

- **Optimization 2: <Input channel reduction: atomics>**

- Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (**2 points**)
- ☐ Shared memory matrix multiplication and input matrix unrolling (**3 points**)
- ☐ Kernel fusion for unrolling and matrix-multiplication (**2 points**)
- ☐ Weight matrix in constant memory (**1 point**)
- ☐ Tuning with restrict and loop unrolling (**3 points**)
- ☐ Sweeping various parameters to find best values (**1 point**)
- ☐ Multiple kernel implementations for different layer sizes (**1 point**)
- ☐ Input channel reduction: tree (**3 point**)
- ☐ Input channel reduction: atomics (**2 point**)

**INPUT CHANNEL REDUCTION: ATOMICS**

- ☐ Fixed point (FP16) arithmetic. (**4 points**)
- ☐ Using Streams to overlap computation with data transfer (**4 points**)
- ☐ An advanced matrix multiplication algorithm (**5 points**)
- ☐ Using Tensor Cores to speed up matrix multiplication (**5 points**)
- ☐ Overlap-Add method for FFT-based convolution (**8 points**)
- ☐ Other optimizations: please explain

<answer here>

- How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?  
<This optimization works by getting rid of the need for a for loop by allowing multiple threads to write to the same global memory location at around the same times using atomic operations. Atomic operations ensure that if any thread at a moment writes a value to the global memory, that value will not accidentally be overwritten by another thread. As a result input channels can do work in parallel, which should increase the speed of the convolution by getting rid of the sequential for loop. This modification would synergies well with something like constant memory, as it remains relatively independent of the constant memory optimization.>

- List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	<0.376371 ms>	<0.941186 ms>	<0m1.279s>	<0.86>
1000	<1.65392 ms>	<7.57432 ms>	<0m10.815s>	<0.886>
10000	<15.7643 ms>	<75.4195 ms>	<1m47.923s>	<0.8714>

- Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of  
<The optimization was not successful in improving performance, as both the total

execution times and op times went up. There is likely some overhead in increasing the size of each block (since we added a third dimension for input channel parallelism). Also, there were a small number of input channels anyway, so the optimization was not very noticeable. >

0% 10 20 30 40 50 60 70 80 90 100%  
 |---|---|---|---|---|---|---|---|---|---|  
 \*\*\*\*\*

Exported successfully to  
 /build/report1.sqlite

Generating CUDA API Statistics...  
 CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum
Name					
77.1	1098639241	8	137329905.1	8135	603996490
cudaMemcpy					
15.3	217419189	8	27177398.6	70188	214037695
cudaMalloc					
6.5	91957447	6	15326241.2	3002	76048638
cudaDeviceSynchronize					
1.0	14245586	6	2374264.3	22817	14103181
cudaLaunchKernel					
0.2	2792429	8	349053.6	69434	974655
cudaFree					

Generating CUDA Kernel Statistics...

Generating CUDA Memory Operation Statistics...  
 CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum
Name					
100.0	91718101	2	45859050.5	15728854	75989247
conv_forward_kernel					
0.0	2848	2	1424.0	1376	1472
do_not_remove_this_kernel					
0.0	2720	2	1360.0	1344	1376

prefn\_marker\_kernel

#### CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum
Name					
-----	-----	-----	-----	-----	-----
93.2	1019488638	2	509744319.0	416403012	603085626 [CUDA memcpy DtoH]
6.8	73805870	6	12300978.3	1504	39059544 [CUDA memcpy HtoD]

#### CUDA Memory Operation Statistics (KiB)

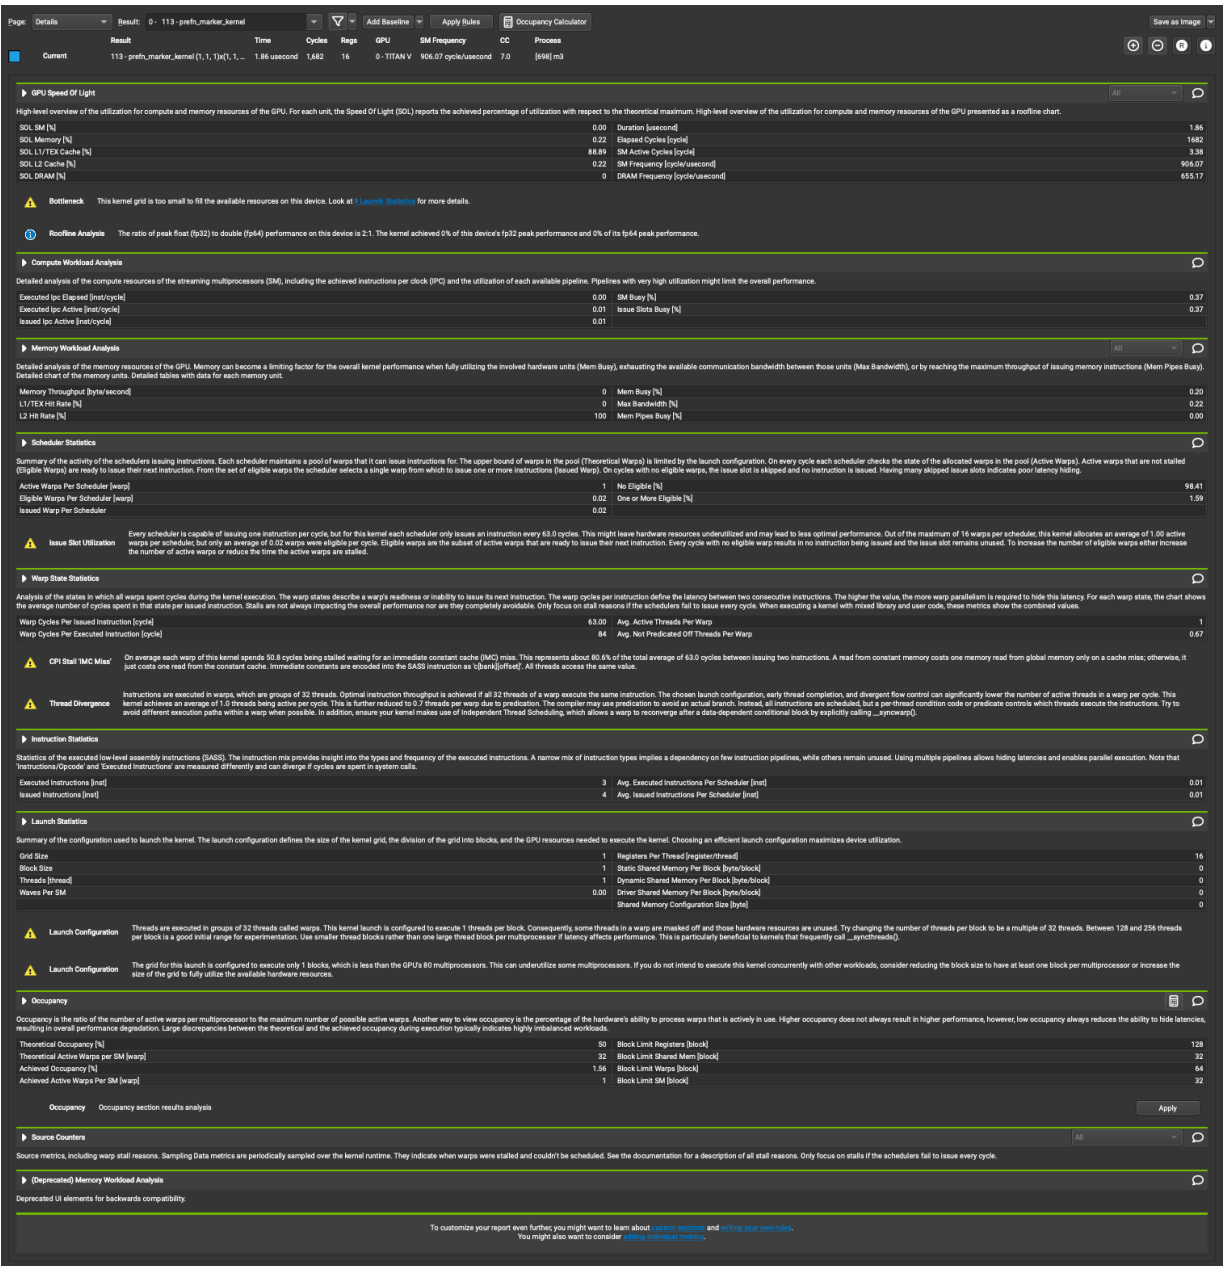
Total	Operations	Average	Minimum	Maximum
Name				
-----	-----	-----	-----	-----
1722500.0	2	861250.0	722500.000	1000000.0 [CUDA memcpy DtoH]
538919.0	6	89819.0	0.004	288906.0 [CUDA memcpy HtoD]

#### Generating Operating System Runtime API Statistics...

#### Operating System Runtime API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum
Name					
-----	-----	-----	-----	-----	-----
33.3	94409383418	958	98548416.9	44928	100314577
sem_timedwait					
33.3	94381346070	957	98622096.2	76138	100395926
poll					
22.0	62333209484	2	31166604742.0	22724892940	39608316544
pthread_cond_wait					
11.3	32008534773	64	500133355.8	500044444	500174930
pthread_cond_timedwait					
0.0	120450582	911	132218.0	1059	20093878
ioctl					
0.0	22826252	9072	2516.1	1220	18704

read					
0.0	18408261	26	708010.0	1128	18334532
fopen					
0.0	3913684	98	39935.6	1155	1650623
mmap					
0.0	1515003	101	15000.0	4472	27967
open64					
0.0	326197	5	65239.4	38423	93773
pthread_create					
0.0	141657	3	47219.0	42867	53421
fgets					
0.0	102195	15	6813.0	2415	14183
write					
0.0	98205	2	49102.5	42746	55459
pthread_mutex_lock					
0.0	80713	17	4747.8	1145	13709
munmap					
0.0	74374	3	24791.3	3043	46997
fopen64					
0.0	58180	39	1491.8	1073	1700
fcntl					
0.0	50180	8	6272.5	1165	10100
fflush					
0.0	30593	5	6118.6	2900	7671
open					
0.0	24572	10	2457.2	1058	9256
fclose					
0.0	14247	2	7123.5	4301	9946
pthread_cond_signal					
0.0	13459	2	6729.5	5736	7723
socket					
0.0	8108	1	8108.0	8108	8108
pipe2					
0.0	7846	1	7846.0	7846	7846
connect					
0.0	3638	2	1819.0	1243	2395
fwrite					
0.0	2023	1	2023.0	2023	2023
bind					
0.0	1032	1	1032.0	1032	1032 listen



• What references did you use when implementing this technique?  
<I used lecture 18, which was histogramming and atomic operations.>

• Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization.

For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling  
<answer here>

```
// Insert your GPU convolution kernel code here
int W_grid = ceil((float)Width_out / TILE_WIDTH);
int m = blockIdx.x;
int h = (blockIdx.y / W_grid) * TILE_WIDTH +
threadIdx.y;
int w = (blockIdx.y % W_grid) * TILE_WIDTH +
threadIdx.x;
int z = blockIdx.z;
float acc = 0.0f;

int c = threadIdx.z;

if(h < Height_out && w < Width_out){
//for(int c = 0; c < Channel; c++){
    for(int p = 0; p < K; p++){
        for(int q = 0; q < K; q++){
            acc += in_4d(z, c, h + p, w + q) *
mask_4d(m, c, p, q);

        }
    }
//}

    atomicAdd(&out_4d(z, m, h, w), acc);
}
```

In the host code a modification was made to add a z dimension to each block, to represent the input channel.

- **Optimization 3: <Weight Matrix in Constant Memory>**

- Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (**2 points**)
- ☐ Shared memory matrix multiplication and input matrix unrolling (**3 points**)
- ☐ Kernel fusion for unrolling and matrix-multiplication (**2 points**)
- ☐ Weight matrix in constant memory (**1 point**)

**WEIGHT MATRIX IN CONSTANT MEMORY**

- ☐ Tuning with restrict and loop unrolling (**3 points**)
- ☐ Sweeping various parameters to find best values (**1 point**)
- ☐ Multiple kernel implementations for different layer sizes (**1 point**)
- ☐ Input channel reduction: tree (**3 point**)
- ☐ Input channel reduction: atomics (**2 point**)
- ☐ Fixed point (FP16) arithmetic. (**4 points**)
- ☐ Using Streams to overlap computation with data transfer (**4 points**)
- ☐ An advanced matrix multiplication algorithm (**5 points**)
- ☐ Using Tensor Cores to speed up matrix multiplication (**5 points**)
- ☐ Overlap-Add method for FFT-based convolution (**8 points**)
- ☐ Other optimizations: please explain

<answer here>

- How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

<Our optimization works by reducing the number of access to global memory by loading the convolution kernels as constant memory, which takes fewer clock cycles to access. I think the optimization will increase performance due to reuse of some elements of every single convolutional filter. I think this becomes Significant for large convolution filter sizes. This optimization would synergize nicely with the atomic operations operations optimization due to its reliance on accessing global memory.>



- List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size  
 Op Time 1  
 Op Time 2  
 Total Execution Time  
 Accuracy  
 100  
 <0.3387 ms>  
 <0.783927 ms>  
 <0m1.198s>  
 <0.86>  
 1000  
 <1.52063 ms>  
 <6.48102 ms>  
 <0m10.422 s>  
 <0.886>  
 10000  
 <14.5867 ms>  
 <64.2765 ms>  
 <1m41.936s>  
 <0.8714>

- Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off  
 <This kernel performed slightly worse than the baseline in op time and total execution time. I suspect the kernel size was not big enough for the advantages of reuse to come into play, and instead I suspect there was a greater overhead from loading the entire convolution kernel from global memory to constant memory.>

Name	Time(%)	Total Time	Calls	Average	Minimum	Maximum
-----	-----	-----	-----	-----	-----	-----
75.8	1064133872	6	177355645.3	18050	575181643	
cudaMemcpy	17.2	241291910	6	40215318.3	283589	237961736
cudaMalloc	5.5	77550874	6	12925145.7	2611	62968824

cudaDeviceSynchronize	1.2	16964279	6	2827379.8	22897	16794763
cudaLaunchKernel	0.2	3385717	8	423214.6	1739	1505472
cudaFree	0.0	37472	2	18736.0	10065	27407
cudaMemcpyToSymbol						

Generating CUDA Kernel Statistics...

Generating CUDA Memory Operation Statistics...  
 CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum
Name					
-----	-----	-----	-----	-----	-----
100.0	77475471	2	38737735.5	14521035	62954436
conv_forward_kernel					
0.0	2656	2	1328.0	1280	1376
prefn_marker_kernel					
0.0	2624	2	1312.0	1248	1376
do_not_remove_this_kernel					

CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum
Name					
-----	-----	-----	-----	-----	-----
92.7	980928707	2	490464353.5	406643722	574284985 [CUDA memcpy DtoH]
7.3	77535661	6	12922610.2	1472	41242162 [CUDA memcpy HtoD]

CUDA Memory Operation Statistics (KiB)

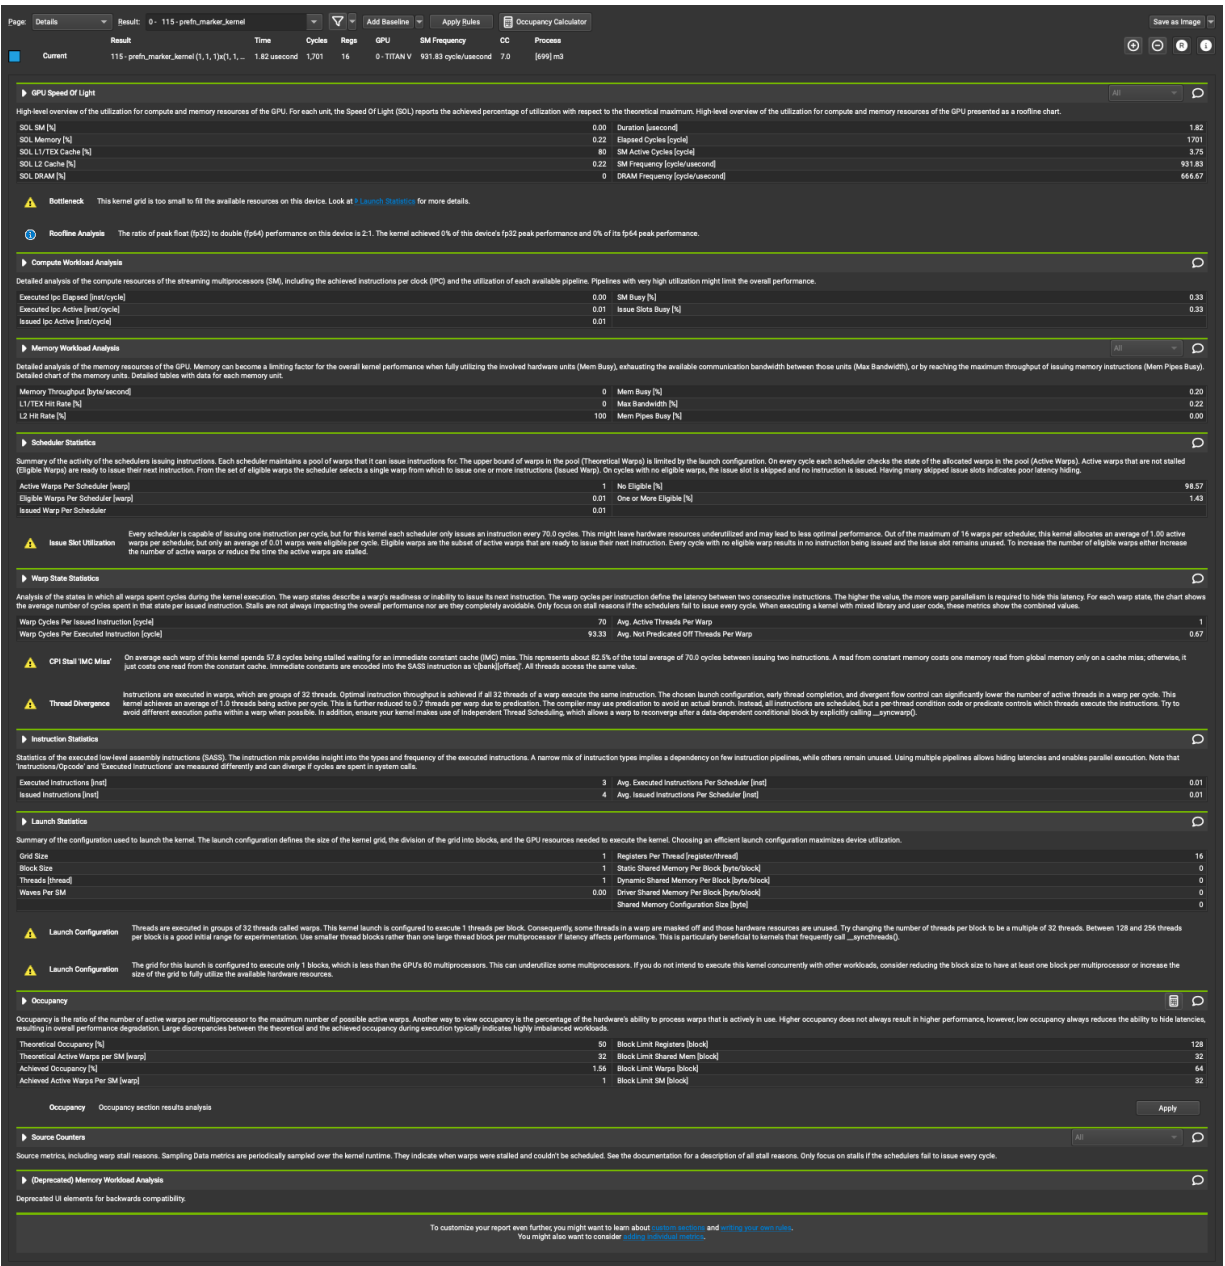
Total	Operations	Average	Minimum	Maximum
Name				
-----	-----	-----	-----	-----
1722500.0	2	861250.0	722500.000	1000000.0 [CUDA memcpy

DtoH]	538919.0	6	89819.0	0.004	288906.0 [CUDA memcpy
HtoD]					

Generating Operating System Runtime API Statistics...  
Operating System Runtime API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum
Name					
33.3	95465670242	968	98621560.2	52360	100356230
sem_timedwait					
33.3	95378665141	967	98633573.1	65796	100365652
poll					
22.1	63418425309	2	31709212654.5	22962143204	40456282105
pthread_cond_wait					
11.2	32006595905	64	500103061.0	500071669	500159032
pthread_cond_timedwait					
0.0	138726493	905	153288.9	1038	18824794
ioctl					
0.0	22186182	9072	2445.6	1223	19315
read					
0.0	19719675	26	758449.0	1291	19641771
fopen					
0.0	3776971	98	38540.5	1059	1573992
mmap					
0.0	1410426	101	13964.6	4201	46988
open64					
0.0	353836	2	176918.0	82371	271465
pthread_mutex_lock					
0.0	306135	5	61227.0	44522	77096
pthread_create					
0.0	158506	3	52835.3	48859	58949
fgets					
0.0	91279	17	5369.4	1152	12959
munmap					
0.0	77109	15	5140.6	2540	12087
write					
0.0	74376	3	24792.0	3355	49946
fopen64					
0.0	52792	7	7541.7	3990	9904
fflush					
0.0	31643	5	6328.6	3149	8684

open					
0.0	28221	12	2351.8	1042	8345
fclose					
0.0	16090	2	8045.0	6678	9412
socket					
0.0	11214	8	1401.8	1025	2086
fcntl					
0.0	8210	1	8210.0	8210	8210
pipe2					
0.0	7982	2	3991.0	3127	4855
pthread_cond_signal					
0.0	7460	1	7460.0	7460	7460
connect					
0.0	6667	4	1666.8	1001	2582
fwrite					
0.0	2060	1	2060.0	2060	2060
bind					
0.0	1006	1	1006.0	1006	1006
putc					



- What references did you use when implementing this technique?

<I used one of the earlier lectures on DRAM bursting, L2 cache, and memory architecture.>

- Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization.

For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling  
<answer here>

```
__constant__ float MASK[4000];

// #include <cmath>
// #include <iostream>
// #include "gpu-new-forward.h"

// #define TILE_WIDTH 16

// __global__ void conv_forward_kernel(float *output,
const float *input, const float *mask, const int Batch,
const int Map_out, const int Channel, const int Height,
const int Width, const int K)
// {
//     /*
//     Modify this function to implement the forward
pass described in Chapter 16.
//     We have added an additional dimension to the
tensors to support an entire mini-batch
//     The goal here is to be correct AND fast.

//     Function paramter definitions:
//     output - output
//     input - input
//     mask - convolution kernel
//     Batch - batch_size (number of images in x)
//     Map_out - number of output feature maps
//     Channel - number of input feature maps
//     Height - input height dimension
//     Width - input width dimension
//     K - kernel height and width (K x K)
//     */

//     const int Height_out = Height - K + 1;
//     const int Width_out = Width - K + 1;
//     //(void)Height_out; // silence declared but
never referenced warning. remove this line when you
start working
```

```

//      //(void)Width_out; // silence declared but never
referenced warning. remove this line when you start
working

//      // We have some nice #defs for you below to
simplify indexing. Feel free to use them, or create
your own.
//      // An example use of these macros:
//      // float a = in_4d(0,0,0,0)
//      // out_4d(0,0,0,0) = a

//      #define out_4d(i3, i2, i1, i0) output[(i3) *
(Map_out * Height_out * Width_out) + (i2) * (Height_out
* Width_out) + (i1) * (Width_out) + i0]
//      #define in_4d(i3, i2, i1, i0) input[(i3) *
(Channel * Height * Width) + (i2) * (Height * Width) +
(i1) * (Width) + i0]
//      #define mask_4d(i3, i2, i1, i0) MASK[(i3) *
(Channel * K * K) + (i2) * (K * K) + (i1) * (K) + i0]

```

- **Optimization 4: <optimization name>**
  - Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose

that optimization technique.

- ☐ Tiled shared memory convolution (**2 points**) THIS ONE
- ☐ Shared memory matrix multiplication and input matrix unrolling (**3 points**)
- ☐ Kernel fusion for unrolling and matrix-multiplication (**2 points**)
- ☐ Weight matrix in constant memory (**1 point**)
- ☐ Tuning with restrict and loop unrolling (**3 points**)
- ☐ Sweeping various parameters to find best values (**1 point**)
- ☐ Multiple kernel implementations for different layer sizes (**1 point**)
- ☐ Input channel reduction: tree (**3 point**)
- ☐ Input channel reduction: atomics (**2 point**)
- ☐ Fixed point (FP16) arithmetic. (**4 points**)
- ☐ Using Streams to overlap computation with data transfer (**4 points**)
- ☐ An advanced matrix multiplication algorithm (**5 points**)
- ☐ Using Tensor Cores to speed up matrix multiplication (**5 points**)
- ☐ Overlap-Add method for FFT-based convolution (**8 points**)
- ☐ Other optimizations: please explain

<answer here>

- How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?  
<your answer here>

- List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size
Op Time 1
Op Time 2
Total Execution Time
Accuracy
100
<op_time>
<op_time>
<exec_time>
<accuracy>
1000
<op_time>
<op_time>
<exec_time>
<accuracy>



10000  
<op\_time>  
<op\_time>  
<exec\_time>  
<accuracy>

- Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of

- What references did you use when implementing this technique?

- Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization. For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling  
<answer here>